

# Deep Reinforcement Learning

## Homework 2

DQN implementation for Super Mario Bros.

Department: ISA PhD Program

Student ID: 112065802

Name: 蔡睿翊

## Preface

In this assignment, I implemented double dueling DQN to play Super Mario Bros. The environment settings, as well as the processes of training and testing, are introduced in detail as follows.

## Training

By following the instructions in the tutorial of Pytorch and modifying the code in the tutorial for performance enhancement, the implementation of training in this assignment includes the following parts.

## Environment Preprocessing

Since there are some segments that are not related to the game states in the frames, we use the preprocessing modules including `SkipFrame`, `GrayScaleObservation`, `ResizeObservation`, and `FrameStack` (the last one is imported from `gym.wrappers`) provided by OpenAI Gym to make the representation of observed states concise.

```

22 # Environment preprocessing
23 class SkipFrame(gym.Wrapper):
24     def __init__(self, env, skip):
25         """Return only every `skip`-th frame"""
26         super().__init__(env)
27         self._skip = skip
28
29     def step(self, action):
30         """Repeat action, and sum reward"""
31         total_reward = 0.0
32         for i in range(self._skip):
33             # Accumulate reward and repeat the same action
34             obs, reward, done, info = self.env.step(action)
35             total_reward += reward
36             if done:
37                 break
38         return obs, total_reward, done, info

```

```

40 class GrayScaleObservation(gym.ObservationWrapper):
41     def __init__(self, env):
42         super().__init__(env)
43         obs_shape = self.observation_space.shape[2:]
44         self.observation_space = Box(low=0, high=255, shape=obs_shape, dtype=np.uint8)
45
46     def permute_orientation(self, observation):
47         # permute [H, W, C] array to [C, H, W] tensor
48         observation = np.transpose(observation, (2, 0, 1))
49         observation = torch.tensor(observation.copy(), dtype=torch.float)
50         return observation
51
52     def observation(self, observation):
53         observation = self.permute_orientation(observation)
54         transform = T.Grayscale()
55         observation = transform(observation)
56         return observation

```

```

58 class ResizeObservation(gym.ObservationWrapper):
59     def __init__(self, env, shape):
60         super().__init__(env)
61         if isinstance(shape, int):
62             self.shape = (shape, shape)
63         else:
64             self.shape = tuple(shape)
65
66         obs_shape = self.shape + self.observation_space.shape[2:]
67         self.observation_space = Box(low=0, high=255, shape=obs_shape, dtype=np.uint8)
68
69     def observation(self, observation):
70         transforms = T.Compose([T.Resize(self.shape, antialias=True), T.Normalize(0, 255)])
71         observation = transforms(observation).squeeze(0)
72         return observation

```

Before the training process starts, we apply the aforementioned wrappers to the environment to make the observed states in 4 channels, with each channel having frames resized as 84 by 84 pixels.

```

368     # Apply Wrappers to environment
369     env = SkipFrame(env, skip=4)
370     env = GrayScaleObservation(env)
371     env = ResizeObservation(env, shape=84)
372     env = FrameStack(env, num_stack=4)

```

## Double Dueling DQN Architecture

The architecture of the dueling double DQN is implemented as follows. The architecture includes online and target networks. Each of the networks is built by a dueling DQN, which consists of 3 convolution layers operated by ReLU, followed by a FC layer for 2 linear layers to predict the value and advantages (for actions in the action space), respectively. Then the two dueling DQNs are applied to construct double dueling DQN as follows.

```

22  # Dueling Double DQN Architecture
23  class D3QN(nn.Module):
24      def __init__(self, input_dim, output_dim):
25          super().__init__()
26          c, h, w = input_dim
27
28          if h != 84:
29              raise ValueError(f"Expecting input height: 84, got: {h}")
30          if w != 84:
31              raise ValueError(f"Expecting input width: 84, got: {w}")
32
33          self.conv1 = nn.Conv2d(c, 4, 3, padding=1)
34          self.conv2 = nn.Conv2d(4, 8, 3, padding=1)
35          self.conv3 = nn.Conv2d(8, 16, 3, padding=1)
36          self.conv4 = nn.Conv2d(16, 16, 3, padding=1)
37          self.conv5 = nn.Conv2d(16, 16, 3, padding=1)
38
39          self.pool = nn.MaxPool2d(2, ceil_mode=True)
40
41          self.fcval = nn.Linear(144, 20)
42          self.fcval2 = nn.Linear(20, 1)
43          self.fcadv = nn.Linear(144, 20)
44          self.fcadv2 = nn.Linear(20, output_dim)
45
46      def forward(self, x):
47          x = self.pool(F.relu(self.conv1(x)))
48          x = self.pool(F.relu(self.conv2(x)))
49          x = self.pool(F.relu(self.conv3(x)))
50          x = self.pool(F.relu(self.conv4(x)))
51          x = self.pool(F.relu(self.conv5(x)))
52
53          x = x.reshape(x.shape[0], -1)
54
55          advantage = F.relu(self.fcadv(x))
56          advantage = self.fcadv2(advantage)
57          advantage = advantage - torch.mean(advantage, dim=-1, keepdim=True)
58
59          value = F.relu(self.fcval(x))
60          value = self.fcval2(value)
61          return value, advantage

```

# Agent

The agent is implemented as follows.

## Parameters

- Memory size: 100000
- Gamma: 0.9
- Batch size: 32
- Exploration decay in epsilon greedy: 0.95
- Learning rate: 0.0025

```
114 # Agent
115 class Agent:
116     def __init__(self, state_dim, action_dim):
117         self.state_dim = state_dim
118         self.action_dim = action_dim
119
120         self.use_cuda = torch.cuda.is_available()
121
122         # Mario's DNN to predict the most optimal action - we implement this in the Learn section
123         self.online_net = D3QN(self.state_dim, self.action_dim).float()
124         self.target_net = D3QN(self.state_dim, self.action_dim).float()
125         if self.use_cuda:
126             self.online_net = self.online_net.to(device='cuda')
127             self.target_net = self.target_net.to(device='cuda')
128
129         self.exploration_rate = 1
130         self.exploration_rate_decay = 0.95
131         self.exploration_rate_min = 0.1
132         self.curr_step = 0
133
134         self.save_every = 5e5 # no. of experiences between saving Mario Net
135         self.memory = deque(maxlen=100000)
136         self.batch_size = 32
137
138         self.gamma = 0.9
139
140         self.optimizer = torch.optim.Adam(self.online_net.parameters(), lr=0.0025)
141         self.loss_fn = torch.nn.SmoothL1Loss()
142
143         self.burnin = 1e4 # min. experiences before training
144         self.learn_every = 3 # no. of experiences between updates to Q_online
145         self.sync_every = 1e4 # no. of experiences between Q target & Q online sync
```

act

This function chooses an action via an epsilon-greedy strategy with an exploration decay.

```

147 def act(self, state):
148     """Given a state, choose an epsilon-greedy action"""
149     # EXPLORE
150     if np.random.rand() < self.exploration_rate:
151         action_idx = np.random.randint(self.action_dim)
152
153     # EXPLOIT
154     else:
155         state = state[0].__array__() if isinstance(state, tuple) else state.__array__()
156         state = torch.tensor(state).cuda() if self.use_cuda else torch.tensor(state).unsqueeze(0)
157         action_values = self.net(state, model="online")
158         action_idx = torch.argmax(action_values, axis=1).item()
159
160     # decrease exploration_rate
161     self.exploration_rate *= self.exploration_rate_decay
162     self.exploration_rate = max(self.exploration_rate_min, self.exploration_rate)
163
164     # increment step
165     self.curr_step += 1
166     return action_idx
167

```

## cache

This function adds the experiences (states, actions, rewards and corresponding next states) to the memory.

```

168 def cache(self, state, next_state, action, reward, done):
169     """Add the experience to memory"""
170     def first_if_tuple(x):
171         return x[0] if isinstance(x, tuple) else x
172     state = first_if_tuple(state).__array__()
173     next_state = first_if_tuple(next_state).__array__()
174
175     state = torch.FloatTensor(state).cuda() if self.use_cuda else torch.FloatTensor(state)
176     next_state = torch.FloatTensor(next_state).cuda() if self.use_cuda else torch.FloatTensor(next_state)
177     action = torch.LongTensor([action]).cuda() if self.use_cuda else torch.LongTensor([action])
178     reward = torch.DoubleTensor([reward]).cuda() if self.use_cuda else torch.DoubleTensor([reward])
179     done = torch.BoolTensor([done]).cuda() if self.use_cuda else torch.BoolTensor([done])
180
181     self.memory.append((state, next_state, action, reward, done))

```

## recall

This function samples a batch of experiences (states, actions, rewards and corresponding next states) from the memory.

```

184 def recall(self):
185     """Sample experiences from memory"""
186     batch = random.sample(self.memory, self.batch_size)
187     state, next_state, action, reward, done = map(torch.stack, zip(*batch))
188     return state, next_state, action.squeeze(), reward.squeeze(), done.squeeze()

```

## learn

This function first samples a batch of experiences (states, actions, rewards and corresponding next states) from the memory with function `recall`, estimates the current Q-value with the value and advantage of the current state predicted by the online network in function `td_estimate`, calculates the Q-value in the next state with the values and advantages predicted by online and target networks in function `td_target`, and updates the online network with SGD in function `update_Q_online` by calculating the TD error between the current Q-value and the Q-value in the next state.

Besides, this function synchronizes the target network and the online network with function `sync_Q_target` per 10000 steps and saves the model with function `save_model` per 500000 steps.

```

190 def td_estimate(self, state, action):
191     current_V, current_A = self.online_net(state)
192     current_Q = current_V + (current_A - current_A.mean(dim=1, keepdim=True)) # q - batch_size * n_actions
193     return current_Q[np.arange(0, self.batch_size), action] # Q_online(s,a)
194
195 @torch.no_grad()
196 def td_target(self, reward, next_state, done):
197     next_state_V1, next_state_A1 = self.online_net(next_state)
198     next_state_V2, next_state_A2 = self.target_net(next_state)
199     next_state_Q1 = next_state_V1 + (next_state_A1 - next_state_A1.mean(dim=1, keepdim=True))
200     next_state_Q2 = next_state_V2 + (next_state_A2 - next_state_A2.mean(dim=1, keepdim=True))
201     best_action = torch.argmax(next_state_Q1, axis=1)
202     next_Q2 = next_state_Q2[np.arange(0, self.batch_size), best_action]
203     return (reward + (1 - done.float()) * self.gamma * next_Q2).float()
204
205 def update_Q_online(self, td_estimate, td_target):
206     loss = self.loss_fn(td_estimate, td_target)
207     self.optimizer.zero_grad()
208     loss.backward()
209     self.optimizer.step()
210     return loss.item()
211
212 def sync_Q_target(self):
213     self.target_net.load_state_dict(self.online_net.state_dict())
214
215 def learn(self):
216     """Update online action value (Q) function with a batch of experiences"""
217     if self.curr_step % self.sync_every == 0:
218         self.sync_Q_target()
219
220     if self.curr_step % self.save_every == 0:
221         self.save_model()
222
223     if self.curr_step < self.burnin:
224         return None, None
225
226     if self.curr_step % self.learn_every != 0:
227         return None, None
228
229     # Sample from memory
230     state, next_state, action, reward, done = self.recall()
231
232     # Get TD Estimate
233     td_est = self.td_estimate(state, action)
234
235     # Get TD Target
236     td_tgt = self.td_target(reward, next_state, done)
237
238     # Backpropagate Loss through Q_online
239     loss = self.update_Q_online(td_est, td_tgt)
240
241     return (td_est.mean().item(), loss)

```

Then the online network model is saved with the function `save_model`. (PS. Thanks to the TA's suggestion, I found the cause of the model's oversize and updated the way of model saving.)

```

243 def save_model(self):
244     save_path = '112065802_hw2_data'
245     torch.save(self.online_net.state_dict(), save_path)
246     print(f"D3QN model saved to {save_path} at step {self.curr_step}")

```

The training process is implemented with 10000 episodes as follows. The metric logger is implemented to monitor the training process (Note the details of the metric logger is omitted)

here since the main focus of this report is to explain the implementation of double DQN algorithm).

```
380     mario = Agent(state_dim=(4, 84, 84), action_dim=env.action_space.n)
381
382     logger = MetricLogger(save_dir)
383
384     episodes = 10000
385     for e in range(episodes):
386
387         state = env.reset()
388
389         # Play the game!
390         while True:
391
392             # Run agent on the state
393             action = mario.act(state)
394
395             # Agent performs action
396             next_state, reward, done, info = env.step(action)
397
398             # Remember
399             mario.cache(state, next_state, action, reward, done)
400
401             # Learn
402             q, loss = mario.learn()
403
404             # Logging
405             logger.log_step(reward, loss, q)
406
407             # Update state
408             state = next_state
409
410             # Check if end of game
411             if done or info["flag_get"]:
412                 break
413
414             logger.log_episode()
415
416         if (e % 20 == 0) or (e == episodes - 1):
417             logger.record(episode=e, epsilon=mario.exploration_rate, step=mario.curr_step)
418             mario.save_model()
```

## Testing

Since the trained model is to be loaded, the agent is implemented in a different way from the training process as follows.

- The observed states are preprocessed manually with functions of OpenCV in the function `act`.
- The best action with the maximum probability predicted by the loaded model is chosen in each iteration.

```

61 # Agent
62 class Agent:
63     def __init__(self):
64         # for local test
65         self.use_cuda = torch.cuda.is_available()
66         self.net = D3QN((4, 84, 84), 12)
67         if self.use_cuda:
68             self.net.load_state_dict(torch.load('112065802_hw2_data')).cpu()
69         else:
70             self.net.load_state_dict(torch.load('112065802_hw2_data'))
71         self.frames = deque(maxlen=4)
72         self.curr_step = 0
73         self.memory = deque(maxlen=100000)
74
75     def act(self, observation):
76         preprocess_obs = cv2.cvtColor(observation, cv2.COLOR_RGB2GRAY)
77         preprocess_obs = cv2.resize(preprocess_obs, (84, 84), interpolation=cv2.INTER_AREA)
78         while len(self.frames) < 3:
79             self.frames.append(preprocess_obs)
80         self.frames.append(preprocess_obs)
81         preprocess_obs = torch.from_numpy(np.array(self.frames) / 255).float().unsqueeze(0)
82         _, action_values = self.net(preprocess_obs)
83         action_idx = torch.argmax(action_values, axis=1).item()
84
85         # increment step
86         self.curr_step += 1
87
88         return action_idx
89
90     def cache(self, state, next_state, action, reward, done):
91         """Add the experience to memory"""
92         def first_if_tuple(x):
93             return x[0] if isinstance(x, tuple) else x
94         state = first_if_tuple(state).__array__()
95         next_state = first_if_tuple(next_state).__array__()
96
97         state = torch.FloatTensor(state.copy())
98         next_state = torch.FloatTensor(next_state.copy())
99         action = torch.LongTensor([action])
100         reward = torch.DoubleTensor([reward])
101         done = torch.BoolTensor([done])
102
103         self.memory.append((state, next_state, action, reward, done))

```

The testing process is implemented with 50 episodes as follows (with each episode iteratively performing the action, cache, and calculation of the cumulative reward). The average reward is calculated after finishing 50 episodes.



```

120 if __name__ == '__main__':
121     env = gym_super_mario_bros.make('SuperMarioBros-v0')
122     env = JoypadSpace(env, COMPLEX_MOVEMENT)
123
124     env.reset()
125     mario = Agent()
126
127     total_reward = 0
128     episodes = 50
129
130     for e in range(episodes):
131         state = env.reset()
132         episode_reward = 0
133         print(f'Episode {e}')
134         while True:
135             # env.render()
136             action = mario.act(state)
137             next_state, reward, done, info = env.step(action)
138             mario.cache(state, next_state, action, reward, done)
139             episode_reward += reward
140
141             state = next_state
142
143             if done or info['flag_get']:
144                 break
145
146         print(f'Episode reward in episode {e}: {episode_reward}')
147         total_reward += episode_reward
148
149     avg_reward = total_reward/50
150     print(f'Average reward: {avg_reward}')

```

## Postscript

Since I cannot use GPU due to the issues of the driver in my lab's servers until I found another resource on 4/10 (Wednesday), and there was a mistake found in the testing code (but fixed at 1am on 4/11) when uploading my trained model to leaderboard, the progress was delayed again.....(I also handed in my homework 1 one day late)

However, to take responsibility for my own homework, I decided to update and resubmit the training/testing code (with Double Dueling DQN implementation and model saving/loading), the saved (online) model (with size reduced to 55 KB), and this report.